

Towards Fortifying the Safety of IoT Systems

Paper #279

ABSTRACT

Today’s IoT systems include event-driven smart applications (apps) that interact with sensors and actuators. A problem specific to IoT systems is that buggy apps, unforeseen bad app interactions, or device failures, can cause unsafe and dangerous physical states (*e.g.*, unlocking the main door when no one is at home). Detection of flaws or vulnerabilities that lead to such states requires a holistic view of installed apps, component devices, how they are configured, and more importantly, how they interact. In this paper, we present IOTSAN, a novel practical system that uses model checking as a building block to reveal “interaction-level” vulnerabilities by identifying events that can lead the system to unsafe states. In building IOTSAN, we design new techniques to alleviate the state explosion associated with model checking IoT systems. IOTSAN also automatically translates smart apps’ source code into a format amenable to model checking. Finally, to better understand the root cause of a detected vulnerability, we also design an attribution mechanism to identify problematic and potentially malicious apps. We evaluate IOTSAN on the Samsung SmartThings platform. From 76 manually configured systems, IOTSAN detects 147 vulnerabilities. We also evaluate IOTSAN with malicious SmartThings apps from a previous effort. IOTSAN not only detects the potential safety violations but is also able to attribute these apps as malicious.

1. INTRODUCTION

A variety of IoT (Internet-of-Things) systems are already widely available on the market. These systems are typically controlled by *event-driven* smart apps that take as input either sensed data, user inputs, or other external triggers (from the Internet) and command one or more actuators towards providing different forms of automation. Examples of sensors include smoke detectors, motion sensors, and contact sensors. Examples of actuators include smart locks, smart power outlets, and door controls. Popular control platforms on which third-party developers can build smart apps that interact wirelessly with these sensors and actuators include Samsung’s SmartThings [63], Apple’s HomeKit [4], and Amazon’s Alexa [3], among others.

While conceivably, IoT is here to stay, current research studies on security/safety of IoT systems are limited in two fronts. First, they focus on *individual components* of IoT systems: there are papers on the security of communication protocols [19, 27, 34, 50, 60], firmware of devices [70, 1, 76, 62, 13, 18], platforms [24, 45], and smart apps [23, 24, 45, 71]. Very few efforts have taken a holistic perspective of *the IoT system*. Second, most current research efforts only focus on securing the cyberspace, and do not address the safety and security of the physical space, which is one of the key obstacles for real-world IoT deployment [53].

Our thesis is that a holistic view of an IoT system is important *i.e.*, the distributed sensors and actuators, and the apps that interact with them need to be considered jointly. While the compromise of an individual component may lead to the compromise of the whole system, certain complex security and safety issues are only revealed when the interactions between components (*e.g.*, a plurality of poorly designed apps) and possible device failures are considered. These latent problems are very real since apps are often developed by third-party vendors without coordination, and are likely to be installed by one or more users (*e.g.*, family members) at different times. Moreover, both legitimate device failures [28, 74, 73, 25] (*e.g.*, from battery depletion) and induced communication failures (*e.g.*, via jamming [59]) can lead to missed interactions between autonomous components, which in turn can cause the entire system to transition into a bad state. These issues are especially dangerous, because bad or missed interactions can be deliberately induced by attackers via spoofing sensors [68, 66], luring users to install malicious apps [45], or jamming sensor reports.

Goals: In this paper, our goal is to build a holistic system which, given an IoT system and a set of default plus user-defined safety properties with regards to both the cyber and physical spaces, (a) finds if components of an IoT system or interactions between components can lead to bad states that violate these properties; and, (b) attributes the detected violations to either benign mis-configurations or potential malicious apps. With re-

gards to (a) we account for cases wherein failed device(s) or interactions can cause a bad state. With regards to (b) we look for repeated instantiations of unsafe states since malicious apps are likely to consistently try to coerce the IoT system into exploitable bad states (*e.g.*, those described in [45]).

To achieve our goal, we need to solve a set of technical challenges. Among these, the key challenge lies in the scope of the analysis: as the number of IoT devices and apps is already large and is only likely to grow in the future [29, 43], physical replication and testing of IoT systems is hard (due to scale). Thus, it is desirable to build a realistic model of the system, which captures the interactions between sensors, apps, and actuators.

Our solution: We achieve our goal by addressing the above and other practical challenges, in a novel framework IOTSAN (for IoT Sanitizer). In brief, IOTSAN uses model checking as a basic building block. Towards alleviating the state space explosion problem associated with model checking [16], we design two optimizations within IOTSAN to (i) only consider apps that interact with each other, and (ii) eliminate unnecessary interleaving that is unlikely to yield useful assessment of unsafe behaviors. We also design an attribution module which flags potentially malicious apps, and attributes other unsafe states to bad design or mis-configuration.

We develop a prototype of IOTSAN based on the SPIN model checker [37] and apply it to the Samsung SmartThings platform. As one of our contributions, we design an automated model generator that translates apps written in Groovy (the programming language of SmartThings apps) into Promela, the modeling language of SPIN. To evaluate IOTSAN, we define 45 safety properties and consider 150 smart apps with 76 configurations. With this setup, IOTSAN discovered 147 violations of 20 safety properties due to app interactions (135 violations) and device failures (12 violations). In an extreme case, 4 smart apps needed to interact to cause the violation, which is extremely difficult to spot manually. We also evaluate our attribution module with 9 malicious apps from [45] that are relevant to our problem scope (*e.g.*, causing bad physical states). IOTSAN attributes all 9 apps to be potentially malicious.

A summary of our contributions is as follows:

- We map the problem of detecting potential safety issues of an IoT system into a model checking problem. We develop novel pre-processing methods to alleviate the state explosion problem in model checking.
- We design IOTSAN to detect safety violations in IoT systems, and develop a prototype that applies to the Samsung SmartThings platform. We develop tools to automatically translate the app source code into Promela. We evaluate IOTSAN with 150 smart apps in SmartThings’ market place and discover 147 pos-

sible safety violations.

- We propose a method to attribute safety violations to either bad apps or mis-configurations. We evaluate this method with 9 known malicious apps and achieve 100% attribution accuracy.

2. BACKGROUND AND SYNOPSIS

Today’s IoT systems [63, 4, 3, 72, 42, 49, 55] typically consist of three major components viz., (i) a hub and the IoT devices it controls, (ii) a platform (can be the hub, a cloud backend, or a combination of both) where smart apps execute, and (iii) a companion mobile app and/or a web-based app to configure and control the system. Without loss of generality, we design IOTSAN assuming this underlying architecture. Therefore, although our prototype is tailored specific to the SmartThings platform given its recent popularity [23, 24, 45, 71], conceptually, IOTSAN could be applied to other IoT platforms as well. We also use the term “IoT system” to refer to those used in smart homes as in recent papers such as [23, 24, 45, 71] for the purposes of the narrative; however, our approach can apply to other application scenarios (*e.g.*, IoT based enterprise deployments or manufacturing systems).

2.1 Samsung SmartThings

Overview: Like the other systems mentioned above, SmartThings has an associated hub and a companion mobile app, that communicate with a cloud backend via the Internet, using the SSL protocol [7]. Developers can create smart apps using the Groovy programming language. The platform and apps interact with devices through *device handlers*; written in Groovy, these are virtual representations of physical devices that expose the devices’ capabilities. To publish a device handler, a developer needs to get a certificate from Samsung. Typically, smart apps and device handlers are executed in the SmartThings cloud backend inside sandboxes.

Programming model: A smart app subscribes to events generated by device handlers (*e.g.*, motion detected) and/or controls some actuators using method calls (*e.g.*, turn on a bulb). Smart apps can also send SMS and make network calls using the SmartThings’ APIs. A smart app can discover and connect to devices, in two ways. Typically, at installation time, the companion app will show a list of supported devices to a user; once the user configures the app, the list of her chosen devices are returned to the app. The second (lesser-known) way is that SmartThings provides APIs that allow apps to query all the devices that are connected to the hub. Besides subscribing to device events, smart apps can also register callbacks for events from external services (*e.g.*, IFTTT [40]) and timers.

Communications: The hub communicates with IoT devices using a protocol such as ZWave or ZigBee. Ex-

periments using the EZSync CC2531 Evaluation Module USB Dongle [41] of Texas Instruments, reveal that the ZigBee implementation in SmartThings supports four (single hop) MAC layer retransmissions. In addition, SmartThings has an application support sublayer that performs 15 end-to-end retransmissions (for a total of 60 retransmissions of a packet). These are in line with ZigBee specifications as also verified in [9, 2, 52, 47]. Thus, typically, it is rare that the system will transition to unsafe states because of benign packet losses.

2.2 Mis-Configuration Problems

Besides malicious apps, mis-configuration is a common cause for safety violations. Specifically, during the installation of a smart app, a user needs to configure the app with sensor(s) and actuator(s). Poor configurations may cause the IoT system to transition to unsafe physical states. There are many common causes for such mis-configurations, *e.g.*, (i) the app’s description is not clear or misleading, (ii) there are too many configuration options, and (iii) normal users often do not have good domain knowledge to clearly understand the behaviors of smart devices and smart apps. To exemplify these issues, we conduct a user study (more details in §10) where we asked 7 student volunteers to configure various apps as they deemed fit. Among these apps, one app is called *Virtual Thermostat* and describes itself as “Control a space heater or window air conditioner (AC) in conjunction with any temperature sensor, like a SmartSense Multi.” Figure 1 shows the inputs needed from a user, which include a temperature measurement sensor (lines 2-5), the power outlets into which the heater or the AC are plugged (lines 6-9), a desired temperature (lines 10-12), etc. Although the developers use the word *or* and the app only expects either a heater or an AC, 5 out of 7 student volunteers thought the app controls *both* a heater and an AC to maintain the desired temperature and mis-configured the app to control both the AC outlet and the heater outlet. To exacerbate the confusion, as the app is expecting the configuration of outlets (`capability.switch`) instead of the actual devices that are plugged into the outlets (*e.g.*, a smart bulb, TV, AC, or heater). Note that SmartThings displays (exports) all available outlets in the IoT system to the user. As a result of volunteer mis-configurations, when the temperature is higher than a predefined threshold, the *Virtual Thermostat* would turn on both the configured outlets (*i.e.*, both the heater and the AC). This violates the following two commonsense properties: (i) a heater is turned on when temperature is above a predefined threshold and (ii) an AC and a heater are both turned on.

2.3 Using Model Checking as a Building Block

The problem of reasoning if and why the IoT system

```

1 preferences {
2   section("Choose a temperature sensor..."){
3     input "sensor", "capability.temperatureMeasurement",
4       title: "Sensor"
5   }
6   section("Select the heater or air conditioner outlet(s)..."){
7     input "outlets", "capability.switch", title: "Outlets",
8       multiple: true
9   }
10  section("Set the desired temperature..."){
11    input "setpoint", "decimal", title: "Set Temp"
12  }
13  section("When there's been movement from (optional, leave blank)
14    \"to not require motion)..."){
15    input "motion", "capability.motionSensor", title: "Motion",
16      required: false
17  }
18  section("Within this number of minutes..."){
19    input "minutes", "number", title: "Minutes", required: false
20  }
21  section("But never go below (or above if A/C) this value with or
22    \"without motion..."){
23    input "emergencySetpoint", "decimal", title: "Emer Temp",
24      required: false
25  }
26  section("Select 'heat' for a heater and 'cool' for an air
27    \"conditioner..."){
28    input "mode", "enum", title: "Heating or cooling?",
29      options: ["heat", "cool"]
30  }
31 }

```

Figure 1: Example of input info needed from users to configure the app *Virtual Thermostat*.

could transition into a bad physical state is challenging because the number of apps and devices is likely to grow in the future and thus, analyzing all possible interactions between them will be hard. Static analysis tools tend to sacrifice completeness for soundness, and thus result in lots of false positives. In contrast, typical dynamic analyses tools verify the properties of a program during execution, but can lead to false negatives.

Model checking is a technique that checks whether a system meets a given specification [44], by systematically exploring the program’s state. In an ideal case, the model checker exhaustively examines all possible system states to verify if there is any violation of specifications relating to safety and/or liveness properties. However, the complexity of modern system software makes this extremely challenging computationally. So in practice, when the goal is to find bugs, a model checker is usually used as a *falsifier* *i.e.*, it explores a portion of the reachable state space and tries to find a computation that violates the specified property. This is sometimes also called bounded model checking [10, 46, 54, 17, 22].

We adopt model checking as a basic building block since: (i) it provides the flexibility towards verifying all the desired properties with linear temporal logic¹, (ii) it provides concrete counter-examples [6, 69] which are very useful in analyzing why and how the bad states occur, (iii) its holistic nature of checking can enable the capture of interactions among multiple apps, and (iv) it is more efficient than exhaustive testing. However, a successful model checker must address the state explosion problem, *i.e.*, the state space could become unwieldy and requires exponential time to explore.

¹ Linear (or linear-time) temporal logic (LTL) is a modal temporal logic with modalities referring to time. LTL is used to specify/verify properties of reactive systems [6].

Model checking can be grouped into two classes: (a) explicit model checking [15] where progress is made one state at a time and, (b) symbolic model checking [51] that examines sets of states in each step. Literature suggests that neither is a clear winner with regards to yielding complete verification within reasonable times in all settings [5, 21, 39].

Unlike symbolic model checkers, explicit model checkers typically support high level modeling languages (*e.g.*, C-like Promela in SPIN [37]). Thus, we use explicit model checking as our basic building block in designing IOTSAN. Specifically, we use the popular SPIN model checker [37] for checking if a given set of safety properties can be possibly violated. Since an IoT system may be composed of a large number of apps and smart devices, we use SPIN’s verification mode with BITSTATE hashing—an approximate technique that stores the hash code of states in a bitfield instead of storing the whole states. Although the BITSTATE hashing technique does not provide a complete verification, empirical results and theoretical analysis have proved its effectiveness in terms of state coverage [38, 12, 14, 36, 35].

3. SCOPE AND THREAT MODEL

In this work, our goal is to detect safety issues (*i.e.*, vulnerabilities) of IoT systems that are exploitable by attackers to cause the system to transition into bad physical states or leak sensitive information. Safety requirements (*i.e.*, definition of bad states and information leakage) can come from both the users and security experts. Examples of bad physical states are (i) the front door is unlocked when no one is at home, and (ii) a heater is turned off when the temperature is below a predefined threshold and people are asleep. Information leakage (in the context of SmartThings) could be: (i) privacy information can be sent out via only message interfaces (*e.g.*, `sendSms`, `sendSmsMessage`, `sendPush`, `sendPushMessage`) but not via network interfaces (*e.g.*, `httpPost`), and (ii) the recipients of `send` methods should match with the configured phone numbers or contacts. We point out that legitimate apps might use `httpPost` to send some control information (*e.g.*, relating to crashes) back to the server. In such cases, we assume that users dictate whether to allow/disallow such operations (based on their privacy preferences).

We consider all devices (hub, sensors, and actuators), the cloud, and the companion app as our trusted computing base (TCB), and do not consider software attacks against them. However, IOTSAN does mitigate physical attacks that can inject event(s) into the system (*e.g.*, by physically increasing the temperature or spoofing the sensors) or maliciously induced device or communication failures (*e.g.*, by jamming [59]). IOTSAN seeks to identify and prevent such events from leading the system into safety violations. However, providing




Figure 2: Chain of events in an IoT system.

targeted solutions to those attack methods (*e.g.*, detecting/preventing spoofing of sensors or jamming mitigation) is out-of-scope.

We also consider potential bad states that can arise due to natural device failures. Note that many users have reported the failures of their ZigBee and Z-Wave IoT devices (*e.g.*, motion sensors, water leak sensors, presence sensors, contact sensors, and garage door openers) in the SmartThings Community [28, 74, 73, 25]. Failures could also result from device batteries running out. We seek to identify if such device failures can cause an IoT system to transition into a bad physical state.

Malicious apps can exploit weaknesses in the configuration and attack other apps by introducing problematic events. We only seek to attribute an app as possibly malicious and leave the confirmation to human experts or other systems.

4. SYSTEM OVERVIEW

Figure 2 illustrates a high level view of the chain of events in an IoT system. In brief, sensors sense the physical world and convert them into events in the cyber world; these events, in turn, are passed onto apps that subscribe to such events. Upon processing the cyber events these apps may output commands to actuators, which then trigger new physical or cyber events. Apps may also directly generate new cyber events. Therefore, a single event could lead to a large sequence of subsequent cyber/physical events, which could translate to a large state space to check.

Figure 3 depicts the overall architecture of our system IOTSAN. It consists of five modules viz., *App Dependency Analyzer*, *Translator*, *Configuration Extractor*, *Model Generator*, and *Output Analyzer*. In designing IOTSAN, we tackle two main challenges: (i) alleviating the state space explosion with model checking [16] for our context, and (ii) the translation of smart apps’ source code to Promela (to facilitate model checking). We address the first problem partially in the *App Dependency Analyzer* and partially in the *Model Generator*. The second problem is handled partially in the *Translator* and partially in the *Model Generator*.

App Dependency Analyzer (§5): This module constructs dependency graphs to capture interactions between event handlers of different apps and identifies handlers that must be jointly analyzed by the model checker. This precludes the unnecessary analysis of unrelated event handlers.

Translator (§6): To the best of our knowledge, there are no model checking systems that readily verify Groovy




Figure 3: IOTSAN architecture overview.

programs. Thus, we build a translator within IOTSAN, that automatically converts Groovy programs into Promela. In doing so, we address the following challenges:

- **Implicit Types.** In Groovy programs, data types of variables and return types of functions are not explicitly declared. To solve this problem, we design an algorithm to infer data types of variables and return types of functions.
- **Built-in Utilities.** Groovy has many built-in utilities, e.g., `find`, `findAll`, `each`, `collect`, `first`, `+` on list types, and `map`. We manually analyzed the behavior of each utility and translated them into corresponding code in Promela.

The output of the *Translator* module is the Promela code of apps' event handlers.

Configuration Extractor (§7): IoT platforms often provide a companion mobile app and/or web-based app to manage/configure the installed smart apps and devices of an IoT system. This module automatically extracts the system's configurations from the manager app.


Model Generator (§8): This module takes Promela code of event handlers, the configuration of the IoT system, and safety properties (both pre-defined and user-defined) as inputs, and creates the Promela model of the system. We use sequential design in modeling the IoT system instead of concurrent design. This approach significantly reduces the problem size since it eliminates unnecessary interleaving that is unlikely to yield useful assessment of unsafe behaviors. The generated model is then checked by SPIN for possible property violations.

Output Analyzer (§9): This module analyzes the verification logs and attributes safety violations to potentially malicious apps, bad designs or mis-configuration. Based on the result, it provides the user, a suggestion to either remove a bad app(s) or change an app(s)'s configuration.


Our work in perspective: IOTSAN can be considered as a service that considers a user's IoT system holistically (i.e., jointly considers apps, devices and configurations) and checks whether a set of user-defined properties hold. In addition to its ability to detect safety violations of the physical space, it also detects information leakage. It can determine if device failures can cause unsafe states and does not need any changes to legacy apps (such changes might either be unaccept-

Table 1: Comparison of IOTSAN and related work.

Feature	FlowFence [24]	ContextIoT [45]	SmartAuth [71]	SIFT [48]	IotSan
Detects physical safety violations		✓	✓	✓	✓
Detects information leakage	✓	✓	✓		✓
Detects violations due to device failures					✓
Performs violation attribution					✓
Accounts for app interactions				✓	✓
No need for app patches					✓



(a) Dependency graph.



(b) Related sets (each box represents a related set).

Figure 4: Example of a dependency graph and its corresponding related sets.

able to developers, or may result in new unforeseen issues). In Table 1 we show the features that IOTSAN offers compared to the most related recent systems. A discussion of related work is deferred to § 12.

5. APP DEPENDENCY ANALYZER

The model checker should not have to check the interactions between event handlers of different apps if they do not interact with each other. To find event handlers that can interact and thereby jointly influence actuator actions, the first module in IOTSAN constructs a *dependency graph* (DG).

Dependency Graph Construction: Each smart app registers one or more *event handlers* that get notified of events to which it has subscribed. An event handler takes one or more input events, and can induce zero or more output events. Input events are (i) explicitly declared in the `subscribe` commands or, (ii) identified via APIs that read states of smart devices, or (iii) indicated by interrupts at specific times defined by `schedule` method calls. Output events are invoked via APIs and change states of smart devices. We enumerate the input and output events of an app using static analysis (details are straightforward thus omitted due to space constraints).

Once the input and output events are identified, we construct a directed DG as follows. Each event handler is denoted by a vertex in the DG. An edge from a vertex u to a vertex v ($u \rightarrow v$) is added if the output events of u overlap with the input events of v . u is then called the *parent* vertex of the *child* vertex v . The vertices

App's Name	Event Handler	Vertex's ID	Input Events	Output Events
Brighten Dark Places	contactOpenHandler	0	contact/open, illuminance/"..."	switch/on
Let There Be Dark!	contactHandler	1	contact/"..."	switch/on, switch/off
Auto Mode Change	presenceHandler	2	presence/"..."	location/mode
Unlock Door	appTouch	3	app/touch	lock/unlocked
	changedLocationMode	4	location/mode	lock/unlocked
Big Turn On	appTouch	5	app/touch	switch/on
	changedLocationMode	6	location/mode	switch/on

(a)		(b)		(c)	
Set	Vertexes	Set	Vertexes	Set	Vertexes
1	0	1	0, 1	1	3
2	1	2	1, 5	2	2, 4
3	3	3	1, 2, 6	3	0, 1
4	5			4	1, 5
5	2, 4			5	1, 2, 6
6	2, 6				

Example: To illustrate, consider the following example. Table 2 summarizes the event handlers and the associated input/output events with a set of sample smart apps. The description of an event is in the format *attribute/event type* (e.g., contact/open means “a contact sensor is open”); empty quotes (“...”) denote “any” event of that type. Given these apps, we show the DG that is built in Figure 4a. For each vertex, the incoming arrows denote input events and the outgoing arrows denote output events. In this figure, vertex 2 has two children viz., vertex 4 and vertex 6; all of the vertices except vertex 2 are leaf vertices.

The initial related sets constructed as above are incomplete. This is because, two vertices u and v may have common output events but the types of these events could be different or what we call *conflicting*. For example, nodes 0 and 1 have conflicting output events viz., switch/off and switch/on. In such cases, the related sets to which u and v belong, must be merged to account for such conflicts. Table 3b shows the related sets of vertices with potential output conflicts in our example. Note here that to check for such output conflicts, we need to examine $O(E^2)$ links in the worst case (given E output edges from the event handlers); our experiments show that such checks are very fast.

Grey boxes are existing modules/packages. The *SPIN Trans* module in *Bandera* is modified.

The architecture of the *SPIN Trans* module is shown in Figure 1. The process starts with *AppS* code, which is processed by the *SmartThings Handler* and *GParser* to generate a *Groovy AST*. This *Groovy AST* is then translated by the *G2J Translator* and a *Translator* into *Java code*. The *Java code* is processed by the *Java Front-end* (consisting of a *Parser* and a *Code Generator*) to generate a *Java AST*. This *Java AST* is then processed by the *Back-ends* (consisting of *SPIN Trans*, *dSPIN Trans*, and *SMV Trans*) to generate a *BIR* model. The *BIR* model is then processed by the *BIR Constructor* to generate a *Jimple* model. The *Jimple* model is then processed by the *Model Generator* to generate a *Promela model of the IoT system*. The *Promela model of the IoT system* is then processed by the *IoTSan* module.

6. TRANSLATOR

We build IOTSAN using the Bandera Tool Set [31, 32], which is a collection of program analysis, transformation, and visualization components designed to apply model-checking to verify Java source code. Bandera generates a program model and specification in the language of one of several existing model-checking tools (including SPIN, dSpin, SMV, JPF). When a model-checker produces an error trail, Bandera renders the error trail at the source code level and allows the user to step through the code along the path of the trail while displaying values of variables and internal states of Java lock objects [31, 32].

While Bandera produces Promela code for SPIN, it does not handle Groovy code. In order to analyze smart apps for SmartThings, we need to convert their code into Java which is challenging for the following reasons. First, since SmartThings added several language features to Groovy to simplify smart app development, the standard Groovy compiler cannot directly process an app’s code and SmartThings’s compiler is not open sourced. Second, Groovy uses dynamic typing [30] (*i.e.*, data types are checked at run-time) but Java is static typed (*i.e.*, data types are explicitly declared and checked at compile-time). Thus, we need to perform type inference during the translation of Groovy into Java. Lastly, Groovy supports many built-in utilities such as list and map that are not supported by Bandera (*i.e.*, Bandera supports only Java’s *array* type).

6

```

1 private onSwitches() {
2     switches + onSwitches
3 }

```

(a) Groovy's code

```

1 private STSwitch[] onSwitches(){
2     STSwitch[] STSwitchArr0;
3     int arrIndex0 = 0;
4     int index3 = 0;
5     while(index3 < TheBigSwitch_switches.length){
6         STSwitch it = TheBigSwitch_switches[index3];
7         STSwitchArr0[arrIndex0] = it;
8         arrIndex0++;
9         index3++;
10    }
11    int index4 = 0;
12    while(index4 < TheBigSwitch_onSwitches.length){
13        STSwitch it = TheBigSwitch_onSwitches[index4];
14        STSwitchArr0[arrIndex0] = it;
15        arrIndex0++;
16        index4++;
17    }
18    return STSwitchArr0;
19 }

```

(b) Corresponding Java's code

Figure 6: Example of translating a Groovy method into the corresponding Java's method.

is translated into a Java class, whose method comprises of a method's header and a block of statements. The translation procedure of a block is straightforward: iterate through the statement list of the input Groovy block, translate each Groovy statement into Java, add the result to a list of Java statements, and build a Java block from the result list. To implement these, we extended the Groovy compiler (*org.codehaus.groovy*) which is then integrated into the Bandera's front-end.

Handling SmartThings' language features: There are several new language syntaxes introduced in SmartThings. Our *SmartThings Handler* parses these new syntaxes and converts them into vanilla Groovy code using specifications based on the domain knowledge of SmartThings. For instance, (as can be seen in in Figure 1) each *input* function defines a global variable (or a class field) of the app. Therefore, we traverse the Groovy's AST of the app and visit all *input* functions to extract all global variables of the app. In addition, apps can use some predefined objects or variables (*e.g.*, *location*) and APIs (*e.g.*, *setLocationMode*), which are not defined in vanilla Groovy. Therefore, we manually add definitions of these global objects. Regarding their method definitions, we leave them as empty during the translation phase (so that the code compiles), but populate them during the modeling phase (see §8).

Type inference: Although the Groovy Compiler *org.codehaus.groovy* already has a sub-package *CompileStatic* for performing static type inference, it only works when the argument type and the return type of a method are given. In other words, a variable declared inside a method can take different runtime types depending on the argument type. Thus, we still need to infer the argument and return type statically. To do so, we consult the calling context of each method invocation by recursively tracking the arguments and return values to their corresponding anchor points—declaration of variables with explicit types (Groovy supports static

typing as well), assignment to constant values (*e.g.*, we can infer that the type of variable *a* is numeric from *def a = 0*), assignment to return values of known APIs, and known objects and their properties. The inference procedure works roughly as follows. When traversing the AST of a method, we store the names and data types of variables at anchor points; the types of other variables are inferred by propagating the types from anchor points. This is done iteratively until we find no more new variables whose type can be inferred.

Handling Groovy's built-in utilities: Another challenge arises when we translate Groovy into Java for use with Bandera. We find that Bandera understands only a very basic set of Java. For instance, it supports only the *array* type natively. In contrast, Groovy's collection types (*e.g.*, *Collection*, *List*, *ArrayList*, *Set*, *Map*, and *HashSet*) all need to be translated into Java's *array* type. We support the popular collection types that are commonly used in smart apps. An example is shown in Figure 6 that translates one Groovy list into a corresponding Java implementation using array. Since the type of *switches* and *onSwitches* is *List of STSwitch*, we infer the return type of *onSwitches()* method as *List of STSwitch*, which is translated into Java's array type (*i.e.*, *STSwitch[]*). The *+* operation on *List* type (line 2 in Figure 6a) is automatically translated into corresponding Java's code (lines 2-17 in Figure 6b). Finally, since this method is a non-void method, we add an explicit *return* statement (line 18 in Figure 6b).

7. CONFIGURATION EXTRACTOR

IoT platforms typically provide a mobile companion app and/or a web-based app to manage and configure smart apps and devices. For Samsung SmartThings, we develop a crawler in Java, using the *Jsoup* package to automatically extract the system's configuration from the management web app [67]. Given a SmartThings account (user's name and password), the crawler logs in to the management web app and extracts (i) installed devices, (ii) installed smart apps, and (iii) configurations of apps. Moreover, whenever a user install a new generic smart device (*e.g.*, a smart power outlet), we have an interface to get the device association info (*e.g.*, this new outlet is used to control an AC) from the user. Extracted configuration is then saved to a file and used later by the *Model Generator*. The process is straightforward and we omit the details in the interest of space.

8. MODEL GENERATOR

Modeling SmartThings in Promela: To analyze the safety of the IoT system, we need to model two key components (not contained in the smart app code): (i) the SmartThings platform and its interactions with smart apps and (ii) IoT devices and their interactions with smart apps. **With regards to the SmartThings**

platform, we first need to model its platform SDKs, which are unfortunately not open sourced. Thus, we model the major classes and methods based on the SmartThings documentation [64]. The example classes include *Command*, *Attribute*, *Event*, *Mode*, *State*, and *Location*. In addition, the SmartThings platform allows apps to register callback functions (*i.e.*, event handlers) via *subscribe* and *schedule* methods. We model these special registration functions so as to invoke callbacks at appropriate times.

We model IoT devices (sensors and actuators) according to their specifications. We point out that both sensors and actuators in SmartThings can generate events of interest to apps. For instance, a motion sensor can generate motion active/inactive events whereas a door lock (actuator) can generate status update events (locked/unlocked). Each device is modeled as having an event queue and a set of notifiers to inform the smart apps that have subscribed to specific types of events. Currently, we support 30 different IoT devices. Note that SmartThings has a built-in “location manager” to manage environmental events (*e.g.*, *sunrise* and *sunset*) and “location modes” (*e.g.*, *Home*, *Away*, and *Night*). It is modeled as an actuator in our model, which can take commands such as users leaving home and may trigger the mode to change from *Home* to *Away*.

Algorithm 1 shows the pseudo code of the main process that models behaviors of a SmartThings system. The model checker enumerates all possible permutations of the input physical events up to a maximum number of events per user’s configuration to exhaustively verify the system. At each iteration, a sensor and a corresponding physical event in the permutation space are selected (line 2). Then, the selected sensor updates its state and event queue, and notifies its subscribers of the state change event (line 3). When an event is pending, it is dispatched to the subscribed apps and the corresponding event handlers of apps are invoked to handle the event (lines 4-6). Each event handler may send some commands to some actuators, which may generate some new cyber events and trigger event handlers of the subscribers.

To model natural or induced (*e.g.*, using jamming [59]) device/communication failures, when generating a sensor event we emulate two scenarios: (i) the sensor is available/online and (ii) the sensor is unavailable/offline. Similarly, whenever receiving a command from a smart app, an actuator may be either available or unavailable. If a device is unavailable, it will not change its state and therefore *not* broadcast a state change event to its subscribers.

Concurrency Model: One can consider two design choices to model the SmartThings system *viz.*, concurrent and sequential design. With the concurrent approach, each device and each smart app is modeled us-

Algorithm 1 Modeling a SmartThings system

```

1: for  $i = 1$  to maximum number of events do {Main event loop of
   an IoT system}
2:   Select a sensor and a corresponding event in the permutation
   space {Generate a physical event}
3:   sensor_state.update(evt)
4:   while any event pending do
5:     dispatch_event(evt) {Dispatch the pending event to
       the subscribed apps and invoke the corresponding
       app_event_handler(evt) to process the event}
6:   end while
7: end for
   {sensor_state.update(evt)}
8: if  $evt \neq$  current state of the sensor then
9:   Add the  $evt$  to the event queue
10:  Update the state of the sensor
11:  Notify the subscribers of the state change event
12: end if
   {app_event_handler(evt)}
13: if some conditions hold then
14:   Send some command to some actuator {Invoke actua-
       tor_state.update(evt), which may subsequently generate some
       new event}
15: end if
   {actuator_state.update(evt)}
16: Verify conflicting and repeated commands violations
17: if  $evt \neq$  current state of the actuator then
18:   Add the  $evt$  to the event queue
19:   Update the state of the actuator
20:   Notify the subscribers of the state change event
21: end if

```

ing a process (*i.e.*, *proctype* in Promela). There is also a process for generating the sensed and environmental events. The processes communicate with each other using message passing (*i.e.*, *chan* in Promela). This modeling option accurately captures the nature of the SmartThings framework. However, it takes a very long verification time due to the interleaving of concurrent processes, causing the state space to explode. In addition, it is unclear if it allows us to discover more unsafe states as IOTSAN already enumerates all possible permutations of events. For instance, it may not matter “when” exactly the next event is generated with respect to a smart app, since the app processes one event at a time anyway. We have also run the two design options with several small systems and found that the sequential approach discovered all of the violations that the concurrent model found. For these reasons, we use a single process for the whole system in a sequential design which uses *inline* methods to model the behavior of devices and smart apps. The devices, smart apps, and event generators, communicate via shared global variables. This significantly mitigates the state space explosion issue.

Safety Properties: We seek to verify 45 properties of the following types:

- *Free of conflicting commands:* When a single external event happens, an actuator should not receive two conflicting commands (*e.g.*, both on and off) – (1 property).
- *Free of repeated commands:* When a single event happens, an actuator should not receive multiple repeated commands of the same type or with the same payload

Table 4: Sample safe physical states.

Category	Number of properties	Sample property
Thermostat, AC, and Heater	5	Temperature should be within a predefined range when people are at home
Lock and door control	8	The main door should be locked when no one is at home
Location mode	3	Location mode should be changed to Away when no one is at home
Security and alarming	14	An alarm should strobe/siren when detecting smoke
Water and sprinkler	3	Soil moisture should be within a predefined range
Others	5	Some devices should not be turned on when no one is at home

```

1 SmartThings0.prom:2690 (state 295) [generatedEvent.evtType = notpresent]
2 SmartThings0.prom:2689 (state 332) [g_STPresSensorArr.element[STPresSensorIndex].subNotifiers[index2] =
  g_STPresSensorArr.element[STPresSensorIndex].subNotifiers[index2] + 1]
3 SmartThings0.prom:2725 (state 757)
  [((g_STPresSensorArr.element[AutoModeChange_people.element[0].gArrIndex].subNotifiers[AutoModeChange_people.element[0].eventCountIndex] > 0))]
4 SmartThings0.prom:2728 (state 759)
  [g_STPresSensorArr.element[AutoModeChange_people.element[0].gArrIndex].subNotifiers[AutoModeChange_people.element[0].eventCountIndex] =
  g_STPresSensorArr.element[AutoModeChange_people.element[0].gArrIndex].subNotifiers[AutoModeChange_people.element[0].eventCountIndex] - 1]
5 SmartThings0.prom:1913 (state 937) [((((location.mode==AutoModeChange_newMode)))]
6 SmartThings0.prom:2308 (state 1797) [ST_Command.evtType = Away]
7 SmartThings0.prom:2438 (state 1765) [location.mode = HandleLocationEvt_mode]
8 SmartThings0.prom:2451 (state 1788) [location.subNotifiers[index0] = location.subNotifiers[index0] + 1]
9 SmartThings0.prom:2784 (state 346) [(((location.subNotifiers[UnlockDoor_location] > 0))]
10 SmartThings0.prom:2787 (state 348) [location.subNotifiers[UnlockDoor_location] = location.subNotifiers[UnlockDoor_location] - 1]
11 SmartThings0.prom:1832 (state 596) [ST_Command.evtType = unlock]
12 SmartThings0.prom:2357 (state 665) [HandleSTLockEvt_state = 48]
13 SmartThings0.prom:2553 (state 783) [g_STLockArr.element[m0_0_JJCTEMP_0.gArrIndex].currentLock = HandleSTLockEvt_state]
14 spin: _spin_nvr.tmp:3, Error: assertion violated
15 spin: text of failed assertion: assert(!(((g_STPresSensorArr.element[alicePresence_STPresSensor].currentPresence != 18)||
  (g_STLockArr.element[doorLock_STLock].currentLock!=48))))

```

Figure 7: Example violation log (filtered).

– (1 property). The latter could indicate a potential DoS or replay attack.

- *Safe physical states*: Table 4 shows some sample safe physical states that the user desires the system to satisfy. These kinds of properties can be verified using linear temporal logic (LTL) [6] – (38 properties). We envision that a more complete list will likely be provided by safety regulations associated with the IoT industry in the future.
- *Free of other known suspicious app behaviors—security-sensitive command and information leakage*: Examples of security-sensitive commands are *unsubscribe* (disabling an app’s functionality) and creating fake events; information leakage can occur with *sendSms* and *httpPost*. When *sendSms* is triggered, for instance, we check whether the recipient matches with the configured phone number to prevent leakage – (4 properties).
- *Robustness to device failure*: An app should quickly check that a command sent to an actuator was acted upon to be robust to device failures. Upon detecting a failure, the app should notify users via SMS/Push messages. This property can be verified using LTL as well – (1 property).

Note that we provide users with an interface to select the list of safety properties they want to verify.

Example: Consider the smart home of a single owner Alice (say), which comprises of a smart lock that controls the main door viz., **Door Lock**, and a presence sensor viz., **Alice’s Presence** (which checks if Alice is at home). Assume that Alice installs two smart apps: *Auto Mode Change*, which manages the location mode based on the events from **Alice’s Presence**; and *Unlock Door*, which unlocks the **Door Lock** based on ex-

plicit user input, or a “location mode” change event. When this system is analyzed by the model checker, a violation is detected as described below.

Figure 7 shows the violation log (filtered) as a counterexample that SPIN created upon verification. The format of each line in the violation log is as follows: file name (*SmartThings0.prom*), line number, state number, and the executed code. In particular, the counterexample has the following steps. **(1)** The event *not present* is generated by **Alice’s presence** if Alice leaves home (line 1) and its subscribers are notified of this state change (line 2). **(2)** The app *Auto Mode Change* reads and processes this state change event (lines 3-5) and notifies the location manager to change the location mode to *Away* (line 6). **(3)** The location manager changes its mode and notifies its subscribers of this change (lines 7-8). **(4)** The app *Unlock Door* reads and processes this mode change event (lines 9-10) and sends an *unlock* command to the device **Door Lock** (line 11), which unlocks the door (lines 12-13). Thus, the system enters an unsafe physical state (*i.e.*, the main door is unlocked when no one is at home) (lines 14-15).

Upon a closer inspection, the description of the *Unlock Door* app suggests that it unlocks the door *only when a user touches the screen on the app*. However, in its implementation, it also unlocks the door whenever the location mode changes (*i.e.*, an inconsistency between the app’s description and its implementation).

9. OUTPUT ANALYZER

The *Output Analyzer* attributes a violation to either a mis-configuration or a malicious app. It uses a heuristic-based algorithm to make this assessment. The algorithm consists of two phases. In the first phase, when a user installs a new smart app, the output analyzer enu-

merates all possible configurations for this new smart app. It verifies if the user-defined properties hold with each configuration independently. If the proportion of violations (violation ratio) is greater than a predefined threshold (*e.g.*, 90%), the new smart app is attributed as a malicious app.

If this is not the case, in the second phase, the new app is verified in conjunction with other apps that were previously installed by the user. Again, all configurations are considered. If the violation ratio is greater than a predefined threshold, the new app is attributed as a bad app and a report is provided to the user. Otherwise, the violation is attributed to mis-configuration and suggestions of safe configurations with regards to the user defined properties are provided. If there is no violation, a successful verification is reported.

10. EVALUATIONS

All our experiments (model checking) are performed on a MacBook Pro with macOS Sierra, 2.9 GHz Intel Core i5, 16 GB 1867 MHz DDR3, and 256 GB SSD. We use the properties discussed in §8 and check if there are violations. We also look at other performance metrics such as the running times, and the scale ratio (which quantifies the reduction in the number of event handlers to be jointly verified) to evaluate IOTSAN.

10.1 Test Cases and Configurations

We perform four different sets of experiments described below. The first three examine the fidelity with which bad apps and configurations are identified. The last set evaluates the performance of different design choices we make.

Market apps with expert configurations: We examine the safety properties with 150 apps (assuming that these are benign) from the Samsung SmartThings’ market place. We (the authors of this paper) came up with reasonable but independent configurations for the apps (based on common sense with regards to how the apps may be used). To illustrate, consider the app *Virtual Thermostat*, the required input to which is shown in Figure 1. Assuming that the following devices are deployed: (1) one temperature sensor (*myTempMeas*), (2) one outlet to control the heater (*myHeaterOutlet*), (3) one outlet to control the air conditioner (*myACOutlet*), (4) one outlet to control the light in the living room (*livRoomBulbOutlet*), (5) one outlet to control the light in the bedroom (*bedRoomBulbOutlet*), (6) one outlet to control the light in the bathroom (*batRoomBulbOutlet*), (7) one motion sensor in the living room (*livRoomMotion*), and (8) one motion sensor in the bathroom (*batRoomMotion*). Our configuration is as follows: *myTempMeas* for the temperature sensor (line 3 in Figure 1), *myACOutlet* for “outlets” (line 7 in Figure 1), 75 as the “setpoint” temperature when people

are present (line 11 in Figure 1), *livRoomMotion* for “motion” (line 15 in Figure 1), 10 “minutes” for turning off the AC/heater when no motion is sensed (line 19 in Figure 1), 85 as the “emergencySetPoint” temperature at which the AC is turned on regardless of whether people are present (line 22 in Figure 1 to set), and “cool” for “mode” (line 28 in Figure 1).

We randomly divide the 150 apps into six groups (25 apps per group) with one configuration each, and feed them into IOTSAN. Upon encountering a violation, we remove the minimum number of the associated apps (*e.g.*, if there are two apps causing conflicting commands, we randomly remove one of them); we then iterate the process. The experiment stops when no violation is detected. These experiments are performed with and without device failures.

Market apps with non-expert configurations:

To eliminate biases, we also conduct a user study where we request 7 independent student volunteers to configure 10 groups of apps with the assumption that they would deploy them at home. Each group comprises of about 5 related apps (as determined by our app dependency analyzer). A group receives one configuration from each volunteer and this leads to a total of 70 configurations. Our Office of Research Integrity determined that there was no need to go through an IRB approval process (since no private information is collected).

Malicious apps: We consider 25 malicious apps created by the effort in [45]. In this set, we find that only 9 apps are relevant to our evaluations (*e.g.*, affect the physical state and can be compiled correctly by the SmartThings’ own web-based IDE). There are four apps that IOTSAN cannot currently handle viz., *Midnight Camera*, *Auto Camera*, *Auto Camera 2*, and *Alarm Manager*, since they dynamically discover and control the devices in the system; we will extend IOTSAN to handle such apps in future work. We evaluate whether IOTSAN correctly attributes these malicious apps when they are installed together with other apps. The configurations of the 9 malicious apps are identical to those in [45]. We also choose 11 potentially bad apps (found via the previous experiments) from the market place to make an input data set of 20 bad apps. In conjunction, we select 10 good apps from the market place to create a reasonable input set. Here, we specifically evaluate the fidelity of our attribution module.

Known bad configurations: We evaluate the performance of concurrent *versus* sequential design. We use two bad groups of apps viz., (Auto Mode Change, Unlock Door) and (Brighten Dark Places, Let There Be Dark), and one good group of apps viz., (Good Night, It’s Too Cold) that control 3 switch devices, 3 motion sensors, and 1 temperature measurement sensor.


10.2 Identifying Unsafe Configurations

Table 5: Verification results with market apps.


Violation type	Number of violations	Example violated property	Apps related to example
Conflicting commands	8	A light receives “on” and “off” simultaneously	(Brighten Dark Places, Let There Be Dark)
Repeated commands	10	A light receives repeated “on” commands	(Automated light, Brighten My Path)
Unsafe physical states	20	A heater is turned off at night when temperature is below a predefined threshold	(Energy Saver)
		The main door is unlocked when people are sleeping at night	(Light Follows Me, Light Off When Close, GoodNight, Unlock Door)

Table 6: Verification results with market apps, with volunteer configuration.

Violation type	Number of violated properties	Number of violations
Conflicting commands	1	19
Repeated commands	1	12
Unsafe physical states	8	66



(a) Example violation due to bad app interactions.



(b) Example violation due to a device failure. Dotted arrows are expected events that do not occur due to the failure of the motion sensor.

Figure 8: Violation examples. The boxes represent apps and high level abstractions are shown for inputs/outputs.

Market apps with expert configurations: Table 5 summarizes the results from our first set of experiments in the absence of device failures. The apps in parenthesis jointly cause a violation. We find 38 violations of 11 properties, some of which can be very dangerous from a user’s perspective. For example, there is violation where “The main door is unlocked when people are sleeping at night”, which involves 4 apps. The interactions between the apps that lead to this violation is shown in Figure 8a: when lights are turned off at night a mode change is initiated by the *Good Night* app, which in turn causes the unsafe action of unlocking the main door by the *Unlock door* app.

Device failures cause violations of 9 additional properties with some dangerous cases. One such case is showcased in Figure 8b. When people leave home, the *Make it so* app should automatically lock the entrance door; however, due to the failure of the motion sensor, the *Make it so* app is not triggered and thus, the door is left unlocked. Moreover, this failure also causes *NO* notification to be sent to law enforcement upon physical intrusion. An alarming discovery is that none of the analyzed apps check if the commands sent to the actuators were actually carried out (which might not be the case if the device has failed).

Market apps with non-expert configurations: We summarize the verification results from the second set of experiments in Table 6. From 10 groups of apps

Table 7: (a) Scalability with dependency graphs. (b) Runtimes with concurrent and sequential design.

(a)				(b)		
Group	Original Size	New Size	Scale Ratio	Number of events	Concurrent	Sequential
1	37	11	3.4	1	1s	1s
2	27	5	5.4	2	56.5s	1s
3	34	23	1.5	3	139m	1s
4	30	12	2.5	4	forever	1s
5	42	19	2.2	5		1s
6	34	6	5.7	6		4.2s
Mean scale ratio				7		16.3s

with 70 configurations (discussed earlier), we find 97 violations of 10 properties. For example, the property “An AC and a heater are both turned on” is violated by 21 configurations across 5 groups. Note that in some configuration multiple properties are violated and thus, the number of violations is more than the number of configurations.

10.3 Violation Attribution

IOTSAN attributes *all* of the ContextIoT’s malicious apps [45] correctly when each is independently considered with violation ratios of 100 % (recall §9). Two apps violated the information leakage property as the command *httpPost* was executed. Two apps violated the “using security-sensitive command property”, *i.e.*, they generated fake carbon monoxide detection events and an *unsubscribe* is executed. The remaining 5 apps violated safety properties in the physical space, *e.g.*, *a main door is unlocked when no one is at home* and, *when smoke is detected, a water valve switch is turned off*. From among the 11 market apps, 6 were detected with a 100% violation ratio, both when verified independently and in conjunction with other apps; they were thus attributed as bad apps. The remaining were attributed to cause violations (with 70% or lower violation ratio) due to bad configurations: in other words, there existed safe configurations with no violations.

10.4 Scalability

Table 7a shows the scalability benefits of our app dependency analyzer in the above experiments with 150 market apps. In this table, “*Original Size*” is the total number of event handlers of a group and “*New Size*” is the number of event handlers of the largest related set after running the *App Dependency Analyzer* module. On average, *App Dependency Analyzer* reduced the problem size by a factor 3.4x.

10.5 Concurrent v.s. Sequential

Model checkers using both concurrent and sequential

design were able to detect all violations very quickly (within a single second). Table 7b shows the runtimes of the two models with a good group of apps (2 apps and 7 devices), which does not violate any property. We can see that sequential design significantly reduces the runtime of the verification. Note that *forever* means the experiment ran for a week and then was forced to stop. Moreover, we also verified the runtime of our sequential approach with a much bigger system, which comprises of 5 related apps and 10 devices and does not have any violation. The verification time for 10 events is about 5 hours, which is quite reasonable for a laptop with limited computing resources.

11. LIMITATIONS

While our prototype of IOTSAN has been shown to be very effective in identifying bad apps and unsafe configurations, it has the following limitations. *First*, the SPIN model checker has a predefined threshold for the size of Promela code (and cannot handle a file size greater than this). Depending on apps' source code sizes and dependencies among the apps, IOTSAN can handle a system with about 30 apps. We assume that users are unlikely to have many more than this today and will investigate further scalability in the future. *Second*, we require smart apps to explicitly subscribe to specific devices they want to control and cannot handle smart apps that dynamically discover devices and interact with them. Such apps are very dangerous since they can control any device without permissions from users. Identifying such apps and ensuring that they do not compromise the physical state is beyond the scope of this effort. *Third*, in Algorithm 1, we let the model checker enumerate all possible permutations of the event types; thus, it may consider scenarios that are unlikely to happen in the real world (*e.g.*, the temperature is set to a minimum value in the first iteration and set to a maximum value in the second one). However, we include these scenarios to catch bad or malicious apps. If such scenarios can be eliminated, the state explosion issue can be further mitigated. *Fourth*, we do not explicitly model the behavior of the physical environment after an actuator executes an command (*e.g.*, the system temperature should increase after a heater is turned on). However, such physical changes are implicitly covered by the way the model checker exhaustively verifies a system.

12. RELATED WORK

IoT Security: Current research on IoT security can be roughly divided into three categories that focus on devices [61, 26, 33], protocols [27, 34, 50, 60], and platforms. There have been efforts addressing information leakage and privacy [11, 75, 65, 8, 77], system-level threats [57], and vulnerabilities of firmware images [18].

Fernandes *et al.*, have recently reported security-critical design flaws in the IoT permission model that could expose smart home users to significant harm such as break-ins [23]. To address these, they propose FlowFence [24], a system that requires smart apps to declare their intended data flow patterns. It then explicitly embeds some extra code into the smart app's structure to block undeclared flows. ContextIoT [45], provides contextual integrity by supporting fine-grained context identification for sensitive actions, to help users perform effective access control. SmartAuth [71] generates a user interface that facilitates educated authorizations based on the app's functions and operations. These efforts propose modifications to a smart app's source code and the platform, to enforce good behaviors of smart apps at run time. In contrast, our work statically identifies possible violation of given physical/cyber safety properties of IoT systems without requiring any app modifications.

Model Checking: Model checking has been used to verify system-level threats of IoT systems. IoTRiskAnalyzer [58] is a probabilistic model checking system that generates system and threat models to capture the risk exposure of each input configuration. IoTSAT [56] utilizes Satisfiability Modulo Theories (SMT) [20] to formally model the generic behavior of IoT systems. SIFT [48] takes app's rules as inputs and uses a synthesis engine to generate code that is specific to the deployment environment. SIFT then uses white-box model checking to verify that synthesized IoT apps do not violate safety policies. In contrast with these efforts, IOTSAN targets developing a practical platform for ensuring the physical safety of today's IoT systems. It not only addresses the practical challenges (*e.g.*, scale issues and making Groovy amenable to model checking) in identifying configurations that violate user properties relating to the physical state, but also addresses robustness (failures) and security issues (malicious app attribution).

13. CONCLUSIONS

Badly designed apps, undesirable interactions between installed apps and/or device failures can cause an IoT system to transition into bad states. In this paper, we design and prototype a framework IOTSAN that uses model checking as a basic building block to identify causes for bad physical/cyber states and provides counter-examples to exemplify these causes. IOTSAN addresses practical challenges such as alleviating state space explosion associated with model checking, and automatic translation of app code into a form amenable for model checking. Our evaluations show that IOTSAN identifies many (sometimes complex) unsafe configurations, and flags considered bad apps with 100% accuracy.

14. REFERENCES

- [1] M. Ahmad. Reliability models for the internet of

- things: A paradigm shift. In *2014 IEEE International Symposium on Software Reliability Engineering Workshops*, pages 52–59, Nov 2014.
- [2] R. Alena, R. Gilstrap, J. Baldwin, T. Stone, and P. Wilson. Fault tolerance in zigbee wireless sensor networks. In *2011 Aerospace Conference*, pages 1–15, March 2011.
- [3] Amazon. Alexa. <https://developer.amazon.com/alexa>, Feb. 2018.
- [4] Apple. Homekit. <https://www.apple.com/ios/home/>, Feb. 2018.
- [5] G. S. Avrunin, J. C. Corbett, and M. B. Dwyer. Benchmarking finite-state verifiers. *International Journal on Software Tools for Technology Transfer*, 2(4):317–320, Mar 2000.
- [6] C. Baier and J. P. Katoen. *Principles of Model Checking*. The MIT Press, Cambridge, Massachusetts, London, England, 2008.
- [7] B. Belleville, P. Biernat, A. Cotenoff, K. Hock, T. Prynn, S. Sankaralingam, T. Sun, and D. Mayer. Internet of things security. <https://www.nccgroup.trust/us/our-research/internet-of-things-security/>, 2018.
- [8] E. Bertino, K.-K. R. Choo, D. Georgakopolous, and S. Nepal. Internet of things (iot): Smart and secure service delivery. *ACM Trans. Internet Technol.*, 16(4):22:1–22:7, Dec. 2016.
- [9] A. Betzler, C. Gomez, I. Demirkol, and J. Paradells. A holistic approach to zigbee performance enhancement for home automation networks. *Sensors*, 14(8):14932–14970, 2014.
- [10] A. Biere, A. Cimatti, E. Clarke, and Y. Zhu. *Symbolic Model Checking without BDDs*, pages 193–207. Springer Berlin Heidelberg, Berlin, Heidelberg, 1999.
- [11] C. Busold, S. Heuser, J. Rios, A.-R. Sadeghi, and N. Asokan. Smart and secure cross-device apps for the internet of advanced things. In *Proc. Financial Cryptography and Data Security*, Puerto Rico, US, 2015.
- [12] T. Cattel. Modelization and verification of a multiprocessor realtime os kernel. In *Proc. 7th FORTE Conference*, pages 35–51, Bern, Switzerland, 1994.
- [13] H. Chandra, E. Anggadajaja, P. S. Wijaya, and E. Gunawan. Internet of things: Over-the-air (ota) firmware update in lightweight mesh network protocol for smart urban development. In *2016 22nd Asia-Pacific Conference on Communications (APCC)*, pages 115–118, Aug 2016.
- [14] J. Chaves. Formal methods at at&t, an industrial usage report. In *Proc. 4th FORTE Conference*, pages 83–90, Sydney, Australia, 1991.
- [15] E. M. Clarke and E. A. Emerson. *Design and synthesis of synchronization skeletons using branching time temporal logic*, pages 52–71. Springer Berlin Heidelberg, Berlin, Heidelberg, 1982.
- [16] E. M. Clarke, W. Klieber, M. Nováček, and P. Zuliani. *Tools for Practical Software Verification*. Springer Berlin Heidelberg, Berlin, Heidelberg, 2012.
- [17] L. Cordeiro, J. Morse, D. Nicole, and B. Fischer. *Context-Bounded Model Checking with ESBMC 1.17*, pages 534–537. Springer Berlin Heidelberg, Berlin, Heidelberg, 2012.
- [18] A. Costin, J. Zaddach, A. Francillon, and D. Balzarotti. A large-scale analysis of the security of embedded firmwares. In *Proc. 23rd USENIX Security Symposium (USENIX Security 14)*, pages 95–110, San Diego, CA, USA, Aug. 2014.
- [19] D. Das and B. Sharma. General survey on security issues on internet of things. *International Journal of Computer Applications*, 139(2), 2016.
- [20] L. De Moura and N. Bjørner. Satisfiability modulo theories: Introduction and applications. *Commun. ACM*, 54(9):69–77, Sept. 2011.
- [21] C. Eisner and D. Peled. *Comparing Symbolic and Explicit Model Checking of a Software System*, pages 230–239. Springer Berlin Heidelberg, Berlin, Heidelberg, 2002.
- [22] R. Ensafi, J. C. Park, D. Kapur, and J. R. Crandall. Idle port scanning and non-interference analysis of network protocol stacks using model checking. In *Proc. of the 19th USENIX conference on security*, pages 17–17, Washington, DC, USA, Aug. 2010.
- [23] E. Fernandes, J. Jung, and A. Prakash. Security analysis of emerging smart home applications. In *Proc. IEEE Symposium on Security and Privacy*, pages 636–654, San Jose, CA, USA, May 2016.
- [24] E. Fernandes, J. Paupore, A. Rahmati, D. Simionato, M. Conti, and A. Prakash. Flowfence: Practical data protection for emerging iot application frameworks. In *Proc. 25th USENIX Security Symposium (USENIX Security 16)*, pages 531–548, Austin, TX, USA, Aug. 2016.
- [25] J. Filippello. Smartsense presence sensor failure. <https://community.smarthings.com/t/smartsense-presence-sensor-failure/16644/9>, Feb. 2018.
- [26] D. Fisher. Pair of bugs open honeywell home controllers up to easy hacks. <https://threatpost.com/pair-of-bugs-open-honeywell-home-controllers-up-to-easy-hacks/113965/>, Feb. 2018.
- [27] B. Fouladi and S. Ghanoun. *Honey, I’m home!!* -

- hacking z-wave home automation systems*. Black Hat, Las Vegas, NV, USA, 2013.
- [28] D. Gray. Devices offline and unavailable. <https://community.smartthings.com/t/devices-offline-and-unavailable/100248>, Feb. 2018.
- [29] A. Groden-Morrison. How the internet of things will drive mobile app development. <https://www.alphasoftware.com/blog/internet-of-things-will-drive-mobile-app-development/>, Feb. 2018.
- [30] Groovy. Type checking extensions. <http://docs.groovy-lang.org/next/html/documentation/type-checking-extensions.html>, Feb. 2018.
- [31] J. Hatcliff and M. Dwyer. *Using the Bandera Tool Set to Model-Check Properties of Concurrent Java Software*, pages 39–58. Springer Berlin Heidelberg, Berlin, Heidelberg, 2001.
- [32] J. Hatcliff and M. Dwyer. About bandera. <http://bandera.projects.cs.ksu.edu/>, Sept. 2018.
- [33] A. Hesseldahl. A hacker’s-eye view of the internet of things. <https://www.recode.net/2015/4/7/11561182/a-hackers-eye-view-of-the-internet-of-things>, Feb. 2018.
- [34] G. Ho, D. Leung, P. Mishra, A. Hosseini, D. Song, and D. Wagner. Smart locks: Lessons for securing commodity internet of things devices. In *Proc. of the 11th ACM on Asia Conference on Computer and Communications Security*, pages 461–472, China, 2016.
- [35] G. J. Holzmann. Proving the value of formal methods. In *Proc. 7th FORTE Conference*, pages 385–396, Bern, Switzerland, 1994.
- [36] G. J. Holzmann. The theory and practice of a formal method: Newcore. In *Proc. 13th IFIP World Computer Congress*, Hamburg, Germany, 1994.
- [37] G. J. Holzmann. The model checker spin. In *IEEE Transaction on Software Engineering*, volume 23, pages 279–295. 1997.
- [38] G. J. Holzmann. An analysis of bitstate hashing. In *Formal Methods in System Design*, volume 13, pages 289–307. 1998.
- [39] G. J. Holzmann. *The Engineering of a Model Checker: the Gnu i-Protocol Case Study Revisited.*, pages 232–244. Springer Berlin Heidelberg, Berlin, Heidelberg, 1999.
- [40] IFTTT. Ifttt homepage. <https://ifttt.com/>, Feb. 2018.
- [41] T. Instruments. Ezsync cc2531 evaluation module usb dongle. <http://www.ti.com/tool/CC2531EMK>, Feb. 2018.
- [42] Intel. Smart buildings. <https://www.intel.com/content/www/us/en/internet-of-things/smart-building-solutions.html>, Feb. 2018.
- [43] B. Intelligence. Here’s how the internet of things will explode by 2020. <http://www.businessinsider.com/iot-ecosystem-internet-of-things-forecasts-and-business-opportunities-2016-2>, Feb. 2018.
- [44] R. Jhala and R. Majumdar. Software model checking. *ACM Computing Surveys (CSUR)*, 41(4):21, 2009.
- [45] Y. J. Jia, Q. A. Chen, S. Wangy, A. Rahmati, E. Fernandes, Z. M. Mao, and A. Prakash. Contextiot: Towards providing contextual integrity to appified iot platforms. In *Proc. Network and Distributed System Security Symposium (NDSS’17)*, San Diego, CA, USA, Mar. 2017.
- [46] D. Kroening and M. Tautschnig. *CBMC – C Bounded Model Checker*, pages 389–391. Springer Berlin Heidelberg, Berlin, Heidelberg, 2014.
- [47] J. S. Lee, Y.-M. Wang, and C. C. Shen. Performance evaluation of zigbee-based sensor networks using empirical measurements. In *2012 IEEE International Conference on Cyber Technology in Automation, Control, and Intelligent Systems (CYBER)*, pages 58–63, May 2012.
- [48] C.-J. M. Liang, B. F. Karlsson, N. D. Lane, F. Zhao, J. Zhang, Z. Pan, Z. Li, and Y. Yu. Sift: Building an internet of safe things. In *Proceedings of the 14th International Conference on Information Processing in Sensor Networks, IPSN ’15*, pages 298–309, New York, NY, USA, 2015. ACM.
- [49] Logitech. Harmony hub. <https://www.logitech.com/en-us/product/harmony-hub>, Feb. 2018.
- [50] N. Lomas. Critical flaw ided in zigbee smart home devices. <https://techcrunch.com/2015/08/07/critical-flaw-ided-in-zigbee-smart-home-devices/>, Feb. 2018.
- [51] K. L. McMillan. *Symbolic Model Checking: An Approach to the State Explosion Problem*. Kluwer Academic Publishers, 1993.
- [52] M. U. Memon, L. X. Zhang, and B. Shaikh. Packet loss ratio evaluation of the impact of interference on zigbee network caused by wi-fi (ieee 802.11b/g) in e-health environment. In *2012 IEEE 14th International Conference on e-Health Networking, Applications and Services (Healthcom)*, pages 462–465, Oct 2012.
- [53] A. Meola. How the internet of things will affect security & privacy.

- <http://www.businessinsider.com/internet-of-things-security-privacy-2016-8>, Feb. 2018.
- [54] F. Merz, S. Falke, and C. Sinz. *LLBMC: Bounded Model Checking of C and C++ Programs Using a Compiler IR*, pages 146–161. Springer Berlin Heidelberg, Berlin, Heidelberg, 2012.
- [55] Microsoft. Azure iot. <https://azure.microsoft.com/en-us/services/iot-hub/>, Feb. 2018.
- [56] M. Mohsin, Z. Anwar, G. Husari, E. Al-Shaer, and M. A. Rahman. Iotsat: A formal framework for security analysis of the internet of things (iot). In *Proc. IEEE Conference on Communications and Network Security (CNS)*, pages 180–188, Philadelphia, PA, USA, Oct. 2016.
- [57] M. Mohsin, Z. Anwar, F. Zaman, and E. Al-Shaer. Iotchecker: A data-driven framework for security analytics of internet of things configurations. *Elsevier Computer and Security*, 70:199–223, Sept. 2017.
- [58] M. Mohsin, M. Sardar, O. Hasan, and Z. Anwar. Iotriskanalyzer: A probabilistic model checking based framework for formal risk analytics of the internet of things. *IEEE Access*, 5:5494–5505, Apr. 2017.
- [59] K. Pelechrinis, M. Iliofotou, and S. V. Krishnamurthy. Denial of service attacks in wireless networks: The case of jammers. *IEEE Communications Surveys Tutorials*, 13(2):245–257, Second 2011.
- [60] E. Ronen, C. O’Flynn, A. Shamir, and A.-O. Weingarten. Iot goes nuclear: Creating a zigbee chain reaction. In *Proc. IEEE Symposium on Security and Privacy*, pages 195–212, San Jose, CA, USA, May 2017.
- [61] E. Ronen and A. Shamir. Extended functionality attacks on iot devices: The case of smart lights. In *Proc. 2016 IEEE European Symposium on Security and Privacy*, pages 3–12, Germany, 2016.
- [62] J. E. G. Sala, R. M. Caporal, E. B. Huerta, J. J. R. Rivas, and J. d. J. Rangel Magdaleno. A smart switch to connect and disconnect electrical devices at home by using internet. *IEEE Latin America Transactions*, 14(4):1575–1581, April 2016.
- [63] Samsung. Smartthings. <https://www.smartthings.com/>, Feb. 2018.
- [64] Samsung. Smartthings’ api documentation. <https://docs.smartthings.com/>, Feb. 2018.
- [65] L. Sha, F. Xiao, W. Chen, and J. Sun. Iiot-sidefender: Detecting and defense against the sensitive information leakage in industry iot. *World Wide Web*, pages 1–30, Apr 2017.
- [66] H. Shin, Y. Son, Y.-S. Park, Y. Kwon, and Y. Kim. Sampling race: Bypassing timing-based analog active sensor spoofing detection on analog-digital systems. In *USENIX Workshop on Offensive Technologies*, 2016.
- [67] SmartThings. Smartthings management page. <https://graph-na02-useast1.api.smartthings.com/>, Feb. 2018.
- [68] Y. Son, H. Shin, D. Kim, Y.-S. Park, J. Noh, K. Choi, J. Choi, Y. Kim, et al. Rocking drones with intentional sound noise on gyroscopic sensors. In *USENIX Security Symposium*, pages 881–896, 2015.
- [69] Spin. What is spin? <http://spinroot.com/spin/whatispin.html>, Feb. 2018.
- [70] A. Tekeoglu and A. S. Tosun. A testbed for security and privacy analysis of iot devices. In *2016 IEEE 13th International Conference on Mobile Ad Hoc and Sensor Systems (MASS)*, pages 343–348, Oct 2016.
- [71] Y. Tian, N. Zhang, Y.-H. Lin, X. Wang, B. Ur, X. Guo, and P. Tague. Smartauth: User-centered authorization for the internet of things. In *26th USENIX Security Symposium (USENIX Security 17)*, pages 361–378, Vancouver, BC, 2017. USENIX Association.
- [72] Vera. Smart home controller. <http://getvera.com/controllers/vera3/>, Feb. 2018.
- [73] A. Viguera. More unavailable devices. <https://community.smartthings.com/t/more-unavailable-devices/98584>, Feb. 2018.
- [74] E. Wilkins. Devices showing up as ‘this device is unavailable at the moment’. <https://community.smartthings.com/t/devices-showing-up-as-this-device-is-unavailable-at-the-moment/94724>, Feb. 2018.
- [75] J. Wilson, D. Boneh, R. S. Wahby, P. Levis, H. Corrigan-Gibbs, and K. Winstein. Trust but verify: Auditing the secure internet of things. In *Proc. of the 15th Annual International Conference on Mobile Systems, Applications, and Services (MobiSys’17)*, pages 464–474, New York, USA, 2017.
- [76] F. Xiao, L. T. Sha, Z. P. Yuan, and R. C. Wang. Vulhunter: A discovery for unknown bugs based on analysis for known patches in industry internet of things. *IEEE Transactions on Emerging Topics in Computing*, PP(99):1–1, 2017.
- [77] Y. Yang, L. Wu, G. Yin, L. Li, and H. Zhao. A survey on security and privacy issues in internet-of-things. *IEEE Internet of Things Journal*, PP:1–10, Apr. 2017.